

Categorical Time-Series Analysis

General Notation

Let (X_t) be a *categorical* time series process, where each X_t takes values in

$$S := \{s_0, \dots, s_m\}, \quad m \in \mathbb{N}.$$

(X_t) may be *nominal*, in which case the order of magnitude of the elements in S is arbitrary and solely used for notational consistency, or *ordinal*, in which case the order of the elements in S reflects the inherent order of the elements.

Furthermore, (X_t) is assumed to be stationary. A distinction is made between two forms of stationarity. A process is said to be weakly stationary if:

$$\begin{aligned} E(X_t) &= \mu_X & \forall t, \\ E(X_t^2) &< \infty & \forall t, \\ \text{Cov}(X_t, X_{t+h}) &= \gamma(h) & \forall t, h. \end{aligned}$$

In contrast, strong stationarity requires that all finite-dimensional distributions are invariant under time shifts, i.e.,

$$F_{X_t, \dots, X_{t+j}} = F_{X_{t+h}, \dots, X_{t+h+j}} \quad \forall t, j, h$$

Marginal Distributions

Regardless of whether we consider a simple categorical process with few categories or attempt to describe a complex, temporally evolving process using moments and measures of serial dependence, the analysis of categorical processes is always built on the examination of the marginal probabilities. However, because ordinal and nominal data encode different types of information, their marginals are typically represented in different ways:

- **Ordinal data:** Most analytical tools for ordinal processes take into account the natural ordering of the elements in S . For this reason, most methods for ordinal data rely on the marginal cumulative distribution function (CDF), which accumulates probabilities in accordance with the natural ordering of the categories:

$$\mathbf{f} = (f_0, \dots, f_{m-1})^\top, \quad f_i := P(X_t \leq s_i).$$

For ordinal data, it is sufficient to consider the cumulative probabilities of the categories $0, \dots, m-1$ since the cumulative probability of the last category m is always equal to 1.

Although one could equivalently estimate the PMF and derive the CDF from it, representing ordinal data through the CDF is often more convenient as it this form that is used to compute metrics for location and dispersion.

- **Nominal data:** Since the categories of nominal variables do not possess a meaningful order, statistical procedures cannot make use of cumulative structure. Therefore the marginal distribution is best represented by the probability mass function (PMF):

$$\mathbf{p} = (p_0, \dots, p_m)^\top, \quad p_i := P(X_t = s_i).$$

Bivariate Probabilities at Lag h

Temporal dependence in categorical time series can be quantified through bivariate lag- h probabilities. These represent the joint probability that the process takes a specific category at time t and another category at time $t + h$

- **Ordinal series (cumulative):**

$$f_{ij}(h) := P(X_t \leq s_i, X_{t+h} \leq s_j), \quad i, j < m.$$

- **Nominal series (joint):**

$$p_{ij}(h) := P(X_t = s_i, X_{t+h} = s_j), \quad i, j \leq m.$$

These quantities serve as input for dependence measures such as the ordinal/ nominal Cohen's K .

Location

For a nominal process, any measure that relies on ordering is inappropriate. Instead, a location measure must depend only on the category probabilities. The mode, defined as the most probable category,

$$\text{Mod}(X_t) = \arg \max_{s_i \in S} P(X_t = s_i),$$

is therefore the natural choice. It provides a meaningful summary of the process by identifying the most frequent category.

For an ordinal process, the natural ordering of the categories allows for additional measures of central tendency. While the mode can still be informative as the most frequent category, the median respects the ordering of categories and is often more useful for describing the central location:

$$\text{Med}(X_t) = \min \left\{ s_i \in S : F(s_i) \geq \frac{1}{2} \right\}.$$

Thus, for ordinal data, both the mode and the median carry information, but they capture slightly different aspects of location: the mode identifies the most typical category, whereas the median indicates the midpoint in the ordered scale.

Measures of Dispersion

The concept of dispersion differs between ordinal and nominal processes. Therefore, two separate measures are used: the *Index of Ordinal Variation* (IOV) for ordinal data and the *Index of Qualitative Variation* (IQV) for nominal data. Both indices are scaled to the interval $[0, 1]$, where a value of 0 represents minimal dispersion and a value of 1 represents maximal dispersion.

$$\text{IOV} = \frac{4}{m} \sum_{i=0}^{m-1} f_i(1 - f_i) \quad \text{IQV} = \frac{m+1}{m} \left(1 - \sum_{i=0}^m p_i^2 \right)$$

Minimal Dispersion

Both measures attain the value 0 if and only if the distribution is degenerate, that is, if all probability mass is concentrated in a single category.

Ordinal case. A one-point distribution assigns all probability mass to a single category $s_j \in S$. For ordinal data, this corresponds to a CDF vector $\mathbf{f}$ of the form

$$\mathbf{f} \in \left\{ \begin{pmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{pmatrix}, \begin{pmatrix} 0 \\ 1 \\ \vdots \\ 1 \end{pmatrix}, \dots, \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{pmatrix} \right\} := \{\mathbf{e}_0, \dots, \mathbf{e}_m\} \subset [0, 1]^{m+1}.$$

In each of these cases, every component f_i satisfies $f_i \in \{0, 1\}$, and thus

$$f_i(1 - f_i) = 0 \quad \forall i.$$

Therefore,

$$\text{IOV} = \frac{4}{m} \sum_{i=0}^{m-1} f_i(1 - f_i) = 0.$$

Conversely, if $\text{IOV} = 0$, then

$$\sum_{i=0}^{m-1} f_i(1 - f_i) = 0,$$

and since each term is nonnegative,

$$f_i(1 - f_i) \stackrel{!}{=} 0 \quad \forall i.$$

This is only possible if $f_i \in \{0, 1\}$ for all i , which implies $\mathbf{f} \in \{\mathbf{e}_0, \dots, \mathbf{e}_m\}$. Hence, $\text{IOV} = 0$ holds *exactly* for one-point distributions.

Nominal case. A degenerate (one-point) distribution assigns all probability mass to a single category $s_j \in S$, which corresponds to a PMF vector of the form

$$\mathbf{p} \in \left\{ \begin{pmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix}, \begin{pmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{pmatrix}, \dots, \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{pmatrix} \right\} := \{\mathbf{e}_0, \dots, \mathbf{e}_m\} \subset [0, 1]^{m+1}.$$

In each of these cases, exactly one component equals 1 and all others equal 0, so

$$\sum_{i=0}^m p_i^2 = 1.$$

Thus,

$$\text{IQV} = \frac{m+1}{m} \left(1 - \sum_{i=0}^m p_i^2 \right) = \frac{m+1}{m} (1 - 1) = 0.$$

Conversely, if $\text{IQV} = 0$, then

$$\frac{m+1}{m} \left(1 - \sum_{i=0}^m p_i^2 \right) = 0 \implies \sum_{i=0}^m p_i^2 = 1.$$

As all probabilities satisfy $0 \leq p_i \leq 1$, it follows that

$$p_i^2 \leq p_i \quad \forall p_i, \text{ and } p_i^2 < p_i \quad \forall p_i \notin \{0, 1\}.$$

Since $\sum_{i=0}^m p_i = 1$, it follows that the only way $\sum_{i=0}^m p_i^2 = 1$ can hold is if exactly one p_i equals 1 and all others are 0. In other words, the distribution is degenerate and $\mathbf{p} \in \{\mathbf{e}_0, \dots, \mathbf{e}_m\}$.

Maximal Dispersion

Ordinal case. Dispersion of an ordinal variable is maximal when half of the probability mass lies in the lowest category and half in the highest category; with all intermediate categories having probability zero. This results in a CDF vector,

$$\mathbf{f} = \left(\frac{1}{2}, \frac{1}{2}, \dots, \frac{1}{2} \right)^\top.$$

Then

$$f_i(1 - f_i) = \frac{1}{2} \left(1 - \frac{1}{2} \right) = \frac{1}{4} \quad \forall i,$$

and thus

$$\text{IOV} = \frac{4}{m} \cdot \sum_{i=0}^{m-1} \frac{1}{4} = \frac{4}{m} \cdot \frac{m}{4} = 1.$$

To show that $\text{IOV} = 1 \implies \mathbf{f} = \left(\frac{1}{2}, \dots, \frac{1}{2} \right)^\top$, recall that the IOV only depends on $\mathbf{f}$ through the sum

$$\sum_{i=0}^{m-1} f_i(1 - f_i).$$

Taking the derivative of $g(f_i) := f_i(1 - f_i)$ and setting it to zero yields:

$$\begin{aligned}\frac{\partial}{\partial f_i} g(f_i) &= 1 - 2f_i, \\ 1 - 2f_i &\stackrel{!}{=} 0, \\ f_i &= \frac{1}{2}.\end{aligned}$$

Since $g(f_i = 0) = g(f_i = 1) = 0$, the function is concave and this critical point is indeed the global maximum. Since it has been already shown that $IOV(\mathbf{f} = (\frac{1}{2}, \dots, \frac{1}{2})^T) = 1$, it directly follows that

$$IOV(\mathbf{f}) = 1 \Leftrightarrow \mathbf{f} = \left(\frac{1}{2}, \dots, \frac{1}{2}\right)^\top$$

Nominal case. For nominal data, maximal dispersion is attained when probability mass is distributed uniformly across all $m + 1$ categories, i.e.

$$p_i = \frac{1}{m + 1}, \quad i = 0, \dots, m.$$

In this case,

$$\sum_{i=0}^m p_i^2 = (m + 1) \frac{1}{(m + 1)^2} = \frac{1}{m + 1},$$

and hence

$$IQV = \frac{m + 1}{m} \left(1 - \frac{1}{m + 1}\right) = 1.$$

Conversely, suppose that $IQV = 1$. Then

$$\frac{m + 1}{m} \left(1 - \sum_{i=0}^m p_i^2\right) = 1 \quad \Rightarrow \quad \sum_{i=0}^m p_i^2 = \frac{1}{m + 1}.$$

Thus, maximizing IQV is equivalent to minimizing

$$f(\mathbf{p}) := \sum_{i=0}^m p_i^2 \quad \text{subject to} \quad \sum_{i=0}^m p_i = 1.$$

Introducing the Lagrangian (with Lagrange multiplier λ)

$$\mathcal{L}(\mathbf{p}, \lambda) := \sum_{i=0}^m p_i^2 - \lambda \left(\sum_{i=0}^m p_i - 1 \right),$$

the first-order conditions are

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial p_j} &= 2p_j - \lambda = 0 \\ \Rightarrow \quad 2p_j &= \lambda, \quad j = 0, \dots, m,\end{aligned}$$

together with the constraint $\sum_{i=0}^m p_i = 1$. Since the same multiplier λ applies to all components, it follows that

$$p_0 = p_1 = \cdots = p_m.$$

Imposing the constraint yields

$$(m+1)p_j = 1 \implies p_j = \frac{1}{m+1} \quad \forall j \in 0, \dots, m.$$

Hence, $f(\mathbf{p})$ is uniquely minimized at

$$\mathbf{p} = \left(\frac{1}{m+1}, \dots, \frac{1}{m+1} \right)^\top,$$

and therefore IQV attains its maximal value at the uniform distribution.

Samples Measures

To estimate the marginal (cumulative) probabilities, we express the relative frequencies through a binarization of the original process (X_t) . For ordinal data, define the indicator vector

$$\mathbf{Z}_t := (Z_{t,0}, \dots, Z_{t,m-1})^\top, \quad Z_{t,i} := \mathbb{I}\{X_t \leq s_i\},$$

and for nominal data analogously

$$\mathbf{Y}_t := (Y_{t,0}, \dots, Y_{t,m})^\top, \quad Y_{t,i} := \mathbb{I}\{X_t = s_i\}.$$

The marginal CDF and PMF vectors are then estimated by the corresponding sample means,

$$\hat{\mathbf{f}} := \frac{1}{n} \sum_{t=1}^n \mathbf{Z}_t, \quad \hat{\mathbf{p}} := \frac{1}{n} \sum_{t=1}^n \mathbf{Y}_t.$$

To derive the asymptotic distribution of $\hat{\mathbf{f}}$ and $\hat{\mathbf{p}}$, a central limit theorem for dependent data is required. Since independence is not assumed, we impose weak dependence conditions on the original process (X_t) .

Specifically, $(X_t)_{t \in \mathbb{Z}}$ is assumed to be strictly stationary and α -mixing with exponentially decaying coefficients.

Following Bradley (2005), let

$$\mathcal{F}_J^L := \sigma(X_t : J \leq t \leq L)$$

denote the σ -field generated by $\{X_t : J \leq t \leq L\}$. The α -mixing coefficients of (X_t) , given that it's strictly stationary, are defined as

$$\alpha_{X_t}(n) := \sup_{A \in \mathcal{F}_{-\infty}^0, B \in \mathcal{F}_n^\infty} |P(A \cap B) - P(A)P(B)|.$$

The process is called α -mixing if $\alpha(n) \rightarrow 0$ as $n \rightarrow \infty$.

We additionally assume exponential decay of the mixing coefficients, i.e., there exist constants $C > 0$ and $\rho \in (0, 1)$ such that

$$\alpha(n) \leq C\rho^n \quad \text{for all } n \in \mathbb{N}.$$

The indicator mappings

$$g_i(X_t) := \mathbb{I}\{X_t \leq s_i\}, \quad h_i(X_t) := \mathbb{I}\{X_t = s_i\},$$

are measurable functions of X_t (see Klenke, 2005, p. 34). Since measurability preserves strict stationarity, the binarized processes $(g_i(X_t))$ and $(h_i(X_t))$ are strictly stationary whenever (X_t) is strictly stationary (see Kallenberg, 2002, p. 558). Consequently, the vector-valued processes $\mathbf{Z}_t$ and $\mathbf{Y}_t$ are strictly stationary as well.

Moreover, each indicator variable $Z_{t,i}$ (or $Y_{t,i}$) is defined on the same probability space as X_t , and the σ -field it generates is a sub- σ -field of $\sigma(X_t)$, i.e.

$$\sigma(Z_{t,i}) \subset \sigma(X_t).$$

The same inclusion holds for time-indexed σ -fields, namely

$$\sigma(Z_{t,i} : J \leq t \leq L) \subset \sigma(X_t : J \leq t \leq L), \quad J \leq L.$$

It therefore follows that, for all $i = 0, \dots, m$ and all $n \in \mathbb{N}$,

$$\alpha_{Z_{t,i}}(n) \leq \alpha_{X_t}(n),$$

and analogously for $Y_{t,i}$. Hence, the vector-valued processes $\mathbf{Z}_t$ and $\mathbf{Y}_t$ are also α -mixing.

The central limit theorem applicable under these assumptions is due to Ibragimov (1962, Theorem 7), which requires strict stationarity and α -mixing. In addition, the theorem assumes the summability condition

$$\sum_{n=1}^{\infty} \alpha(n)^{\lambda/(2+\lambda)} < \infty \quad \text{for some } \lambda > 0.$$

This condition is satisfied under exponentially decaying mixing coefficients. Indeed, if $\alpha(n) \leq C\rho^n$ with constants $C > 0$ and $\rho \in (0, 1)$, then

$$\begin{aligned} \sum_{n=1}^{\infty} \alpha(n)^{\lambda/(2+\lambda)} &\leq C^{\lambda/(2+\lambda)} \sum_{n=1}^{\infty} \rho^{n\lambda/(2+\lambda)} \\ &= C^{\lambda/(2+\lambda)} \frac{\rho^{\lambda/(2+\lambda)}}{1 - \rho^{\lambda/(2+\lambda)}} < \infty, \end{aligned}$$

since $0 < \rho^{\lambda/(2+\lambda)} < 1$ for all $\lambda > 0$.

Moreover, the CLT additionally requires the existence of moments of order $2 + \lambda$ for

some $\lambda > 0$. This condition is trivially satisfied in our setting, since $\mathbf{Z}_t$ and $\mathbf{Y}_t$ are vectors of indicator variables. In particular, for any

$$|Z_{t,i}|^{2+\lambda} \in \{0, 1\},$$

and hence all higher order moments exist. Consequently, for all $i \in 0, \dots, m$ and $\lambda > 0$,

$$\mathbb{E}|Z_{t,i}|^{2+\lambda} < \infty,$$

and analogously for $Y_{t,i}$.

According to the Cramér-Wold device, $\mathbf{Z}_T$ converges in distribution to a multivariate normal distribution $\mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ if and only if for every fixed non-zero vector $\mathbf{a} \in \mathbb{R}^{m+1} \setminus \mathbf{0}$, the scalar projections

$$U_t := \mathbf{a}^\top \mathbf{Z}_t = \sum_{i=0}^m a_i Z_{t,i}$$

converge in distribution to a univariate normal distribution $\mathcal{N}(\mathbf{a}^\top \boldsymbol{\mu}, \mathbf{a}^\top \boldsymbol{\Sigma} \mathbf{a})$.

Since each $\mathbf{Z}_t$ is the binarization vector of X_t , it can only take values in the finite set

$$\{\mathbf{c}_0, \dots, \mathbf{c}_m\}$$

where

$$\mathbf{c}_j := (1, \dots, 1, 0, \dots, 0)^\top$$

has exactly $j + 1$ ones followed by $m - j$ zeros. Consequently, for any fixed $\mathbf{a}$, the projection U_t takes values in the finite set

$$\mathcal{A}_a := \{\mathbf{a}^\top \mathbf{c}_j \mid j = 0, \dots, m\} = \left\{ \sum_{i=0}^j a_i \mid j = 0, \dots, m \right\}$$

Equivalently, U_t can be directly expressed as a deterministic function $h : S \rightarrow \mathcal{A}_a$ of the original state s_i of X_t :

$$h_a(s_j) := \sum_{i=0}^j a_i \quad \forall j = 0, \dots, m,$$

so that $U_t = h_a(X_t)$. Since S is finite and equipped with the discrete σ -algebra 2^S (the power set), any function from S to $\mathcal{A}_a \subset \mathbb{R}$ is measurable. In particular, for any subset $E \subset \mathcal{A}_a$ the preimage

$$h^{-1}(E) := \{s_i \in S : h_a(s_i) \in E\}$$

is a subset of S and therefor belongs to 2^S .

Since U_t is a measurable function of X_t , it remains its properties of strict stationarity and α -mixing with exponentially decaying coefficients.

Accordingly, the summability condition is satisfied (follows from the α -mixing with exponentially decaying coefficients) and since U_t is in the discrete, finite set $\mathcal{A}_a$, there exists

$$M := \max_{j=0, \dots, m} \left| \sum_{i=0}^j a_i \right| < \infty$$

Since $|U_t| \leq M$, we have $|U_t|^{2+\lambda} \leq M^{2+\lambda}$ for every $\lambda > 0$, and therefore

$$E|U_t|^{2+\lambda} < |M|^{2+\lambda} < \infty \quad \forall \lambda > 0$$

Since U_t is strictly stationary, α -mixing and satisfying the summability aswell as the moment condition, the CLT from Ibragimov is applicable. It follows that the sample mean

$$\bar{U}_t := \frac{1}{n} \sum_{t=1}^n U_t$$

satisfies the CLT:

$$\sqrt{n}(\bar{U}_t - \mu_{U_t}) \xrightarrow[n \rightarrow \infty]{d} \mathcal{N}(0, \sigma_a),$$

where

$$\begin{aligned} \mu_{U_t} &= E[U_t] \quad \text{and} \quad \sigma_a = \text{Var}(U_t) + 2 \sum_{h=1}^{\infty} \text{Cov}(U_t, U_{t+h}) \\ &= \sum_{i,j} a_i a_j \left[\text{Cov}(Z_{t,i}, Z_{t,j}) + 2 \sum_{h=1}^{\infty} \text{Cov}(Z_{t,i}, Z_{t+h,j}) \right] \end{aligned}$$

Hence, according to the Cramer-Wold device, the sample mean

$$\bar{\mathbf{Z}}_n := \frac{1}{n} \sum_{t=1}^n \mathbf{Z}_t = \hat{\mathbf{f}}$$

satisfies the following multivariate CLT:

$$\sqrt{n}(\bar{\mathbf{Z}}_n - \mathbf{f}) \xrightarrow[n \rightarrow \infty]{d} \mathcal{N}(\mathbf{0}, \Sigma_{\mathbf{Z}}), \quad \text{with } \Sigma_{\mathbf{Z}} = (\sigma_{\mathbf{Z};i,j})_{i,j=0,\dots,m-1}$$

where $\Sigma_{\mathbf{Z}}$ must satisfy that for all $\mathbf{a} \in \mathbb{R}^m$

$$\mathbf{a}^T \Sigma_{\mathbf{Z}} \mathbf{a} = \sigma_a = \sum_{i,j} a_i a_j \left[\text{Cov}(Z_{t,i}, Z_{t,j}) + 2 \sum_{h=1}^{\infty} \text{Cov}(Z_{t,i}, Z_{t+h,j}) \right],$$

which implies

$$\sigma_{\mathbf{Z};i,j} = \text{Cov}(Z_{t,i}, Z_{t,j}) + 2 \sum_{h=1}^{\infty} \text{Cov}(Z_{t,i}, Z_{t+h,j}).$$

Since

$$\text{Cov}(Z_{t,i}, Z_{t,j}) = E[Z_{t,i} Z_{t,j}] - E[Z_{t,i}] E[Z_{t,j}] = f_{\min\{i,j\}} - f_i f_j$$

and

$$\text{Cov}(Z_{t,i}, Z_{t+h,j}) = E[Z_{t,i} Z_{t+h,j}] - E[Z_{t,i}] E[Z_{t+h,j}] = f_{ij}(h) - f_i f_j$$

we can also write the elements of $\Sigma_{\mathbf{Z}}$ as

$$\sigma_{\mathbf{Z};i,j} = (f_{\min\{i,j\}} - f_i f_j) + 2 \sum_{h=1}^{\infty} [f_{ij}(h) - f_i f_j]$$

According to the multivariate delta-method a random variable $G^* := g^*(\Theta^*)$, where $\Theta^* \in \mathbb{R}^m$, $\Theta^* \stackrel{asy.}{\sim} \mathcal{N}(\Theta, \Sigma^*)$ and $g : \mathbb{R}^m \rightarrow \mathbb{R}$ is a continously diffentiable function, satisfies

$$G^* \stackrel{asy.}{\sim} \mathcal{N}(g^*(\Theta), n^{-1} \mathbf{j} \Sigma \mathbf{j}^\top)$$

where $\mathbf{j}$ is the one-dimensional Jakobi-matrix (i.e. vector) and where $j_i = \frac{\partial g^*(\Theta)}{\partial \theta_i}$.

Since we've shown that $\hat{\mathbf{f}}$ is asymptotically normal we can use this to derive the asymptotic properties of sample measures like the IOV. The respective jacobian is given by:

$$\mathbf{j}_{IOV} = \left(\frac{\partial IOV(\mathbf{f})}{\partial f_0}, \dots, \frac{\partial IOV(\mathbf{f})}{\partial f_{m-1}} \right) = \frac{4}{m} ((1 - 2f_0), \dots, (1 - 2f_{m-1}))$$

which leads to an asymptotic variance of the IOV:

$$\begin{aligned} Var(\widehat{IOV(\mathbf{f})}) &= \mathbf{j}_{IOV} \Sigma_{\mathbf{Z}} \mathbf{j}_{IOV}^\top = \sum_{i=0}^{m-1} \sum_{j=0}^{m-1} \frac{4(1 - 2f_i)}{m} \frac{4(1 - 2f_j)}{m} \sigma_{\mathbf{Z};i,j} \\ &= \frac{16}{m^2} \sum_{i,j=0}^{m-1} (1 - 2f_i)(1 - 2f_j) \sigma_{\mathbf{Z};i,j} \end{aligned}$$

and since $\hat{\mathbf{f}}$ has $\mathbf{f}$ as its asymptotic expectation, it follows

$$E[\widehat{IOV(\mathbf{f})}] \stackrel{asy.}{=} IOV(E[\hat{\mathbf{f}}]) = IOV(\mathbf{f}).$$

In a similiar way the asymptotic properties of other dispersion metrics such as the skew can be analyzed. A simple measure for the skew of an ordinal process is given by:

$$\text{skew}(\mathbf{f}) = \frac{2}{m} \sum_{i=0}^{m-1} f_i - 1.$$

The jacobian for the skew is given by

$$\mathbf{j}_{skew} = \left(\frac{\partial skew(\mathbf{f})}{\partial f_0}, \dots, \frac{\partial skew(\mathbf{f})}{\partial f_{m-1}} \right) = \left(\frac{2}{m}, \dots, \frac{2}{m} \right)$$

such that we end up with asymptotic expectation and variance:

$$\begin{aligned} Var(\widehat{\text{skew}(\mathbf{f})}) &= n^{-1} \mathbf{j}_{skew} \Sigma_{\mathbf{Z}} \mathbf{j}_{skew}^\top = \frac{1}{n} \frac{4}{m^2} \sum_0^{m-1} \sigma_{\mathbf{Z};ij} \\ E(\widehat{\text{skew}(\mathbf{f})}) &= \text{skew}(E(\hat{\mathbf{f}})) = \text{skew}(\mathbf{f}) \end{aligned}$$

The derivations for other dispersion metrics are analogous.

Modelling Missingness

To model the missingness, the concept of **amplitude modulation** is employed. This approach, inspired by signal processing, models a stochastic process $(O_t)_t$ that runs concurrently with the main process $(X_t)_t$ and modulates its "amplitude" in a binary manner:

$$O_t \in \{0, 1\} \quad \forall t$$

where $O_t = 1$ indicates that X_t is fully observed (amplitude = 1), and $O_t = 0$ indicates that X_t is missing (amplitude = 0).

The amplitude modulating process O_t can be:

- **Deterministic:** For each t , $O_t = o_t$ where $o_t \in \{0, 1\}$ is a known, fixed value.
- **Stochastic:** O_t follows a probability distribution with $\pi := P(O_t = 1)$, possibly dependent on the history of X_t , O_t , or external factors.

The amplitude modulation is applied to the vector of binarizations $\mathbf{Z}_t$ to get $(O_t \mathbf{Z}_t)$. We then define the naive estimator

$$\hat{\mathbf{f}} := \frac{1}{n} \sum_{t=1}^n O_t \mathbf{Z}_t$$